

Project Title :Predicting Default Risk of Lending Club Loans

Objective:

Our main objective for this project is to implement some of the classification algorithms in order to build models to classify a good or bad loan with help of selected variables.

Dataset:

The dataset includes detailed information for every loan issued by Lending Club for year 2015.

The dataset from Lending Club contains 74 features that will be employed to train our model for prediction. Not all of the fields are intuitively useful for our learning models, such as the loan ID and the month the last payment was received, and thus we removed such fields.

We also removed fields for which greater than 80% of the values were missing.

Number of Features: 74 Number of records: 403697

Questions:

1) Which features contributes the most?

Feature Selection

1) Lasso Regression 2) Principle Component Analysis 3) Recursive Feature Elimination 4) Random Forest for Feature Ranking

Models

1) Random Forest 2) Logistic Regression 3) Support Vector Machine 4) K Nearest Neighbors(KNN)

```
In [49]: import pandas as pd  
import numpy as np
```

```
In [50]: loans_2015 = pd.read_csv('loan2015new.csv', low_memory=False)
```

```
In [51]: loans_2015.shape
```

```
Out[51]: (403697, 74)
```

```
In [52]: loans_2015.columns
```

```
Out[52]: Index(['id', 'member_id', 'loan_amnt', 'funded_amnt', 'funded_amnt_inv',  
              'term', 'int_rate', 'installment', 'grade', 'sub_grade', 'emp_title',  
              'emp_length', 'home_ownership', 'annual_inc', 'verification_status',  
              'issue_d', 'pymnt_plan', 'url', 'desc', 'purpose', 'title', 'zip_code',  
              'addr_state', 'dti', 'delinq_2yrs', 'earliest_cr_line',  
              'inq_last_6mths', 'mths_since_last_delinq', 'mths_since_last_record',  
              'open_acc', 'pub_rec', 'revol_bal', 'revol_util', 'total_acc',  
              'initial_list_status', 'out_prncp', 'out_prncp_inv', 'total_pymnt',  
              'total_pymnt_inv', 'total_rec_prncp', 'total_rec_int',  
              'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',  
              'last_pymnt_d', 'last_pymnt_amnt', 'next_pymnt_d', 'last_credit_pull_d',  
              'collections_12_mths_ex_med', 'mths_since_last_major_derog',  
              'policy_code', 'application_type', 'annual_inc_joint', 'dti_joint',  
              'verification_status_joint', 'acc_now_delinq', 'tot_coll_amt',  
              'tot_cur_bal', 'open_acc_6m', 'open_il_6m', 'open_il_12m',  
              'open_il_24m', 'mths_since_rcnt_il', 'total_bal_il', 'il_util',  
              'open_rv_12m', 'open_rv_24m', 'max_bal_bc', 'all_util',  
              'total_rev_hi_lim', 'inq_fi', 'total_cu_tl', 'inq_last_12m',  
              'loan_status'],  
              dtype='object')
```

```
In [53]: eighty_count = len(loans_2015)*4 / 5
```

```
In [54]: loans_2015 = loans_2015.dropna(thresh=eighty_count,axis=1)
```

```
In [55]: loans_2015.shape
```

```
Out[55]: (403697, 53)
```

```
In [56]: data_dictionary = pd.read_csv('LCDataDictionary.csv')
```

```
In [57]: data_dictionary.head()  
data_dictionary = data_dictionary.rename(columns={'LoanStatNew': 'name', 'Descript
```

```
In [58]: loan_df = pd.DataFrame(data=loans_2015)
```

```
In [59]: loans_df_dtypes = pd.DataFrame(loan_df.dtypes, columns=['dtypes'])  
loans_df_dtypes = loans_df_dtypes.reset_index()  
loans_df_dtypes['name'] = loans_df_dtypes['index']  
loans_df_dtypes = loans_df_dtypes[['name', 'dtypes']]  
loans_df_dtypes['first value'] = loan_df.loc[0].values  
preview = loans_df_dtypes.merge(data_dictionary, on='name', how='left')
```

In [60]: preview

Out[60]:

| | name | dtypes | first value | descript |
|----|---------------------|---------|--------------|--|
| 0 | id | int64 | 68587652 | A unique assigned ID for loan list |
| 1 | member_id | int64 | 73477494 | A unique assigned Id for borrower meml |
| 2 | loan_amnt | int64 | 25000 | The listed amoun the loan applied by |
| 3 | funded_amnt | int64 | 25000 | The total amc committed to 1 loan at th |
| 4 | funded_amnt_inv | int64 | 25000 | The total amc committec investors for 1 |
| 5 | term | object | 36 months | The numbe payments on loan. Values ar |
| 6 | int_rate | float64 | 5.32 | Interest Rate on l |
| 7 | installment | float64 | 752.87 | The mon payment owec the borrower if 1 |
| 8 | grade | object | A | LC assigned l gr; |
| 9 | sub_grade | object | A1 | LC assigned l subgr; |
| 10 | emp_title | object | Director | The job supplied by Borrower when a |
| 11 | emp_length | object | 1 year | Employment len in years. Poss values ; |
| 12 | home_ownership | object | MORTGAGE | The hc ownership sta provided by bo |
| 13 | annual_inc | float64 | 150000 | The self-repor annual incc provided by 1 |
| 14 | verification_status | object | Not Verified | N |
| 15 | issue_d | object | 15-Dec | The month wr the loan was funi |
| 16 | pymnt_plan | object | n | Indicates payment plan been put in |

| | name | dtypes | first value | descript |
|----|---------------------|---------|---|--|
| 17 | url | object | https://www.lendingclub.com/browse/loanDetail.... | URL for the page with lis d; |
| 18 | purpose | object | credit_card | A category provi by the borrower the |
| 19 | title | object | Credit card refinancing | The loan provided by borro |
| 20 | zip_code | object | 054xx | The first 3 numb of the zip co provided |
| 21 | addr_state | object | VT | The state provi by the borrowe the loa |
| 22 | dti | float64 | 9.54 | A ratio calcula using the borrow tote |
| 23 | delinq_2yrs | int64 | 0 | The number of : days past-d incidences c |
| 24 | earliest_cr_line | object | Feb-96 | The month borrower's earl reported cr |
| 25 | inq_last_6mths | int64 | 0 | The numbe inquiries in pa months (ex |
| 26 | open_acc | int64 | 7 | The number of o credit lines in borrow |
| 27 | pub_rec | int64 | 0 | Numbe derogatory pu recc |
| 28 | revol_bal | int64 | 19339 | Total credit revol bala |
| 29 | revol_util | float64 | 42.5 | Revolving utilization rate the amou |
| 30 | total_acc | int64 | 18 | The total numbe credit lines curre ii |
| 31 | initial_list_status | object | w | The initial lis status of the lo Poss |
| 32 | out_prncp | float64 | 24358 | Remair outstanc principal for t amc |
| 33 | out_prncp_inv | float64 | 24358 | Remair outstanc principal for por i |

| | name | dtypes | first value | descript |
|----|----------------------------|---------|-------------|--|
| 34 | total_pymnt | float64 | 682.67 | Payments received to date for total amount funded |
| 35 | total_pymnt_inv | float64 | 682.67 | Payments received to date for portion of total |
| 36 | total_rec_prncp | float64 | 642.03 | Principal received |
| 37 | total_rec_int | float64 | 40.64 | Interest received |
| 38 | total_rec_late_fee | float64 | 0 | Late fees received |
| 39 | recoveries | float64 | 0 | post charge gross recoveries |
| 40 | collection_recovery_fee | float64 | 0 | post charge collection |
| 41 | last_pymnt_d | object | 16-Jan | Last month payment was received |
| 42 | last_pymnt_amnt | float64 | 701.14 | Last total payment amount received |
| 43 | next_pymnt_d | object | 16-Feb | Next scheduled payment date |
| 44 | last_credit_pull_d | object | 16-Jan | The most recent month LC pulled credit for total |
| 45 | collections_12_mths_ex_med | int64 | 0 | Number of collections in 12 months excluding current |
| 46 | policy_code | int64 | 1 | publicly available policy_code=1\nnr product |
| 47 | application_type | object | INDIVIDUAL | Indicates whether the loan is individual or business |
| 48 | acc_now_delinq | int64 | 0 | The number of accounts on which the borrower is delinquent |
| 49 | tot_coll_amt | int64 | 0 | Total collection amounts ever over the life of the loan |
| 50 | tot_cur_bal | int64 | 430856 | Total current balance of all accounts |
| 51 | total_rev_hi_lim | int64 | 45500 | Maximum total revenue limit |
| 52 | loan_status | object | Current | Current status of the loan |

```
In [61]: preview[:19]
```

```
Out[61]:
```

| | name | dtypes | first value | description |
|----|---------------------|---------|---|---|
| 0 | id | int64 | 68587652 | A unique LC assigned ID for the loan listing. |
| 1 | member_id | int64 | 73477494 | A unique LC assigned Id for the borrower member. |
| 2 | loan_amnt | int64 | 25000 | The listed amount of the loan applied for by t... |
| 3 | funded_amnt | int64 | 25000 | The total amount committed to that loan at tha... |
| 4 | funded_amnt_inv | int64 | 25000 | The total amount committed by investors for th... |
| 5 | term | object | 36 months | The number of payments on the loan. Values are... |
| 6 | int_rate | float64 | 5.32 | Interest Rate on the loan |
| 7 | installment | float64 | 752.87 | The monthly payment owed by the borrower if th... |
| 8 | grade | object | A | LC assigned loan grade |
| 9 | sub_grade | object | A1 | LC assigned loan subgrade |
| 10 | emp_title | object | Director | The job title supplied by the Borrower when ap... |
| 11 | emp_length | object | 1 year | Employment length in years. Possible values ar... |
| 12 | home_ownership | object | MORTGAGE | The home ownership status provided by the borr... |
| 13 | annual_inc | float64 | 150000 | The self-reported annual income provided by th... |
| 14 | verification_status | object | Not Verified | NaN |
| 15 | issue_d | object | 15-Dec | The month which the loan was funded |
| 16 | pymnt_plan | object | n | Indicates if a payment plan has been put in pl... |
| 17 | url | object | https://www.lendingclub.com/browse/loanDetail.... | URL for the LC page with listing data. |
| 18 | purpose | object | credit_card | A category provided by the borrower for the lo... |

```
In [62]: drop_list = ['id','member_id','funded_amnt','funded_amnt_inv',
                    'int_rate','sub_grade','emp_title','issue_d','url']
loan_df = loan_df.drop(drop_list,axis=1)
```

```
In [63]: loan_df.shape
```

```
Out[63]: (403697, 44)
```

```
In [64]: preview[19:38]
```

```
Out[64]:
```

| | name | dtypes | first value | description |
|----|---------------------|---------|-------------------------|---|
| 19 | title | object | Credit card refinancing | The loan title provided by the borrower |
| 20 | zip_code | object | 054xx | The first 3 numbers of the zip code provided b... |
| 21 | addr_state | object | VT | The state provided by the borrower in the loan... |
| 22 | dti | float64 | 9.54 | A ratio calculated using the borrower's total ... |
| 23 | delinq_2yrs | int64 | 0 | The number of 30+ days past-due incidences of ... |
| 24 | earliest_cr_line | object | Feb-96 | The month the borrower's earliest reported cre... |
| 25 | inq_last_6mths | int64 | 0 | The number of inquiries in past 6 months (excl... |
| 26 | open_acc | int64 | 7 | The number of open credit lines in the borrowe... |
| 27 | pub_rec | int64 | 0 | Number of derogatory public records |
| 28 | revol_bal | int64 | 19339 | Total credit revolving balance |
| 29 | revol_util | float64 | 42.5 | Revolving line utilization rate, or the amount... |
| 30 | total_acc | int64 | 18 | The total number of credit lines currently in ... |
| 31 | initial_list_status | object | w | The initial listing status of the loan. Possib... |
| 32 | out_prncp | float64 | 24358 | Remaining outstanding principal for total amou... |
| 33 | out_prncp_inv | float64 | 24358 | Remaining outstanding principal for portion of... |
| 34 | total_pymnt | float64 | 682.67 | Payments received to date for total amount funded |
| 35 | total_pymnt_inv | float64 | 682.67 | Payments received to date for portion of total... |
| 36 | total_rec_prncp | float64 | 642.03 | Principal received to date |
| 37 | total_rec_int | float64 | 40.64 | Interest received to date |

```
In [65]: drop_cols = ['zip_code', 'out_prncp', 'out_prncp_inv',
                    'total_pymnt', 'total_pymnt_inv']
loan_df = loan_df.drop(drop_cols, axis=1)
```

```
In [66]: loan_df.shape
```

```
Out[66]: (403697, 39)
```

```
In [67]: preview[38:]
```

```
Out[67]:
```

| | name | dtypes | first value | description |
|----|----------------------------|---------|-------------|---|
| 38 | total_rec_late_fee | float64 | 0 | Late fees received to date |
| 39 | recoveries | float64 | 0 | post charge off gross recovery |
| 40 | collection_recovery_fee | float64 | 0 | post charge off collection fee |
| 41 | last_pymnt_d | object | 16-Jan | Last month payment was received |
| 42 | last_pymnt_amnt | float64 | 701.14 | Last total payment amount received |
| 43 | next_pymnt_d | object | 16-Feb | Next scheduled payment date |
| 44 | last_credit_pull_d | object | 16-Jan | The most recent month LC pulled credit for thi... |
| 45 | collections_12_mths_ex_med | int64 | 0 | Number of collections in 12 months excluding m... |
| 46 | policy_code | int64 | 1 | publicly available policy_code=1\nnew products... |
| 47 | application_type | object | INDIVIDUAL | Indicates whether the loan is an individual ap... |
| 48 | acc_now_delinq | int64 | 0 | The number of accounts on which the borrower i... |
| 49 | tot_coll_amt | int64 | 0 | Total collection amounts ever owed |
| 50 | tot_cur_bal | int64 | 430856 | Total current balance of all accounts |
| 51 | total_rev_hi_lim | int64 | 45500 | NaN |
| 52 | loan_status | object | Current | Current status of the loan |

```
In [68]: drop_cols = ['total_rec_prncp', 'total_rec_int', 'total_rec_late_fee',
                    'recoveries', 'collection_recovery_fee', 'last_pymnt_d']

loan_df = loan_df.drop(drop_cols, axis=1)
```

```
In [69]: loan_df.shape
```

```
Out[69]: (403697, 33)
```

```
In [70]: loan_df["loan_status"].value_counts()
```

```
Out[70]: Current      377553
Fully Paid    22984
Charged Off   2773
Default       387
Name: loan_status, dtype: int64
```

```
In [71]: loan_df = loan_df[(loan_df["loan_status"] == "Fully Paid") |
                        (loan_df["loan_status"] == "Charged Off")]

mapping_dictionary = {"loan_status":{ "Fully Paid": 1, "Charged Off": 0}}
loan_df = loan_df.replace(mapping_dictionary)
```

```
In [73]: loan_df = loan_df.loc[:, loan_df.apply(pd.Series.nunique) != 1]
```



```
In [74]: for col in loan_df.columns:
          if (len(loan_df[col].unique()) < 4):
              print(loan_df[col].value_counts())
              print()
```

```
36 months    18102
60 months    7655
Name: term, dtype: int64
```

```
MORTGAGE    12824
RENT         10026
OWN          2907
Name: home_ownership, dtype: int64
```

```
Source Verified    10678
Verified           8132
Not Verified       6947
Name: verification_status, dtype: int64
```

```
w    13615
f    12142
Name: initial_list_status, dtype: int64
```

```
Series([], Name: next_pymnt_d, dtype: int64)
```

```
INDIVIDUAL    25756
JOINT          1
Name: application_type, dtype: int64
```

```
0    25612
1     135
2      10
Name: acc_now_delinq, dtype: int64
```

```
1    22984
0     2773
Name: loan_status, dtype: int64
```

```
In [75]: print(loan_df.shape[1])
loan_df = loan_df.drop('application_type', axis=1)
print("We've been able to reduced the features to => {}".format(loan_df.shape[1]))
```

```
31
We've been able to reduced the features to => 30
```

```
In [76]: loan_df = loan_df[(loan_df["initial_list_status"] == "w") |
                           (loan_df["initial_list_status"] == "f")]

mapping_dictionary = {"initial_list_status":{ "w": 1, "f": 0}}
loan_df = loan_df.replace(mapping_dictionary)
```

```
In [77]: loan_df.to_csv("filtered_loans_2015.csv",index=False)
```

Handle missing values and categorical features before feeding the data into a machine learning algorithm

```
In [78]: filtered_loans = pd.read_csv('filtered_loans_2015.csv')
print(filtered_loans.shape)
filtered_loans.head()
```

(25757, 30)

Out[78]:

| | loan_amnt | term | installment | grade | emp_length | home_ownership | annual_inc | verification_st |
|---|-----------|-----------|-------------|-------|------------|----------------|------------|-----------------|
| 0 | 19800 | 36 months | 666.00 | C | 7 years | MORTGAGE | 78924.0 | Not Ver |
| 1 | 35000 | 36 months | 1177.27 | C | 10+ years | RENT | 95000.0 | Source Ver |
| 2 | 20000 | 36 months | 672.73 | C | 6 years | MORTGAGE | 56000.0 | Not Ver |
| 3 | 28000 | 60 months | 635.37 | C | 5 years | MORTGAGE | 96000.0 | Not Ver |
| 4 | 23100 | 60 months | 605.35 | E | 10+ years | MORTGAGE | 55000.0 | Source Ver |

5 rows × 30 columns

```
In [79]: null_counts = filtered_loans.isnull().sum()
print("Number of null values in each column:\n{}".format(null_counts))
```

Number of null values in each column:

| | |
|----------------------------|-------|
| loan_amnt | 0 |
| term | 0 |
| installment | 0 |
| grade | 0 |
| emp_length | 0 |
| home_ownership | 0 |
| annual_inc | 0 |
| verification_status | 0 |
| purpose | 0 |
| title | 0 |
| addr_state | 0 |
| dti | 0 |
| delinq_2yrs | 0 |
| earliest_cr_line | 0 |
| inq_last_6mths | 0 |
| open_acc | 0 |
| pub_rec | 0 |
| revol_bal | 0 |
| revol_util | 13 |
| total_acc | 0 |
| initial_list_status | 0 |
| last_pymnt_amnt | 0 |
| next_pymnt_d | 25757 |
| last_credit_pull_d | 0 |
| collections_12_mths_ex_med | 0 |
| acc_now_delinq | 0 |
| tot_coll_amt | 0 |
| tot_cur_bal | 0 |
| total_rev_hi_lim | 0 |
| loan_status | 0 |
| dtype: | int64 |

Remove entire column next_pymnt_d column as data is missing for entire filtered loans dataset. Drop the missing rows for revol_util

```
In [80]: filtered_loans = filtered_loans.drop("next_pymnt_d",axis=1)
filtered_loans = filtered_loans.dropna()
```

```
In [81]: filtered_loans.shape
```

```
Out[81]: (25744, 29)
```

Convert categorical values to numerical

```
In [82]: print("Data types and their frequency\n{}".format(filtered_loans.dtypes.value_counts()))
```

```
Data types and their frequency
int64      14
object     10
float64     5
dtype: int64
```

```
In [83]: object_columns_df = filtered_loans.select_dtypes(include=['object'])
print(object_columns_df.iloc[0])
```

```
term                36 months
grade                C
emp_length           7 years
home_ownership       MORTGAGE
verification_status  Not Verified
purpose              debt_consolidation
title                Debt consolidation
addr_state           MD
earliest_cr_line     5-Oct
last_credit_pull_d   16-Jan
Name: 0, dtype: object
```

```
In [84]: cols = ['home_ownership', 'grade', 'verification_status', 'emp_length', 'term', 'addr_state']
for name in cols:
    print(name, ':')
    print(object_columns_df[name].value_counts(), '\n')
```

```
home_ownership :
MORTGAGE      12816
RENT           10022
OWN            2906
Name: home_ownership, dtype: int64
```

```
grade :
C       7576
B       5945
D       4663
A       3314
E       2906
F       1048
G        292
Name: grade, dtype: int64
```

```
verification_status :
Source Verified    10671
Verified           8128
Not Verified       6945
Name: verification_status, dtype: int64
```

```
emp_length :
10+ years      8850
2 years        2321
< 1 year       2095
3 years        1952
1 year         1694
4 years        1504
5 years        1460
8 years        1311
n/a            1224
7 years        1199
6 years        1092
9 years        1042
Name: emp_length, dtype: int64
```

```
term :
36 months      18092
60 months       7652
Name: term, dtype: int64
```

```
addr_state :
CA      4153
TX      2068
NY      1862
FL      1772
IL       890
NJ       842
GA       824
PA       808
```

| | |
|----|-----|
| OH | 800 |
| VA | 795 |
| NC | 763 |
| MI | 715 |
| AZ | 666 |
| CO | 659 |
| WA | 623 |
| MD | 620 |
| MA | 526 |
| MN | 512 |
| NV | 430 |
| IN | 423 |
| MO | 367 |
| OR | 362 |
| TN | 362 |
| WI | 347 |
| AL | 327 |
| LA | 312 |
| CT | 289 |
| SC | 272 |
| UT | 247 |
| OK | 223 |
| KY | 223 |
| KS | 200 |
| HI | 172 |
| AR | 164 |
| NM | 151 |
| MS | 139 |
| RI | 119 |
| WV | 119 |
| NH | 105 |
| DE | 87 |
| MT | 73 |
| DC | 66 |
| WY | 60 |
| AK | 57 |
| SD | 56 |
| VT | 46 |
| NE | 31 |
| ME | 9 |
| ND | 8 |

Name: addr_state, dtype: int64

```
In [85]: for name in ['purpose','title']:
          print("Unique Values in column: {}".format(name))
          print(filtered_loans[name].value_counts(),'\n')
```

Unique Values in column: purpose

| | |
|--------------------|-------|
| debt_consolidation | 16080 |
| credit_card | 4526 |
| home_improvement | 1725 |
| other | 1442 |
| major_purchase | 545 |
| medical | 275 |
| small_business | 260 |
| car | 255 |
| moving | 231 |
| house | 203 |
| vacation | 185 |
| renewable_energy | 16 |
| wedding | 1 |

Name: purpose, dtype: int64

Unique Values in column: title

| | |
|-------------------------|-------|
| Debt consolidation | 16072 |
| Credit card refinancing | 4535 |
| Home improvement | 1726 |
| Other | 1442 |
| Major purchase | 543 |
| Medical expenses | 275 |
| Business | 259 |
| Car financing | 255 |
| Moving and relocation | 231 |
| Home buying | 204 |
| Vacation | 185 |
| Green loan | 16 |
| Credit Card/Auto Repair | 1 |

Name: title, dtype: int64

The purpose and title columns do contain overlapping information, but the purpose column contains fewer discrete values and is cleaner, so we'll keep it and drop title. The addr_state column, however, contains too many unique values, so it's better to drop this

```
In [86]: drop_cols = ['last_credit_pull_d','addr_state','title','earliest_cr_line','purpose']
          filtered_loans = filtered_loans.drop(drop_cols,axis=1)
```

Convert Ordinal values to numerical values

```
In [87]: filtered_loans['grade'] = filtered_loans['grade'].map({'A':7,'B':6,'C':5,'D':4,'E':3})
filtered_loans["home_ownership"] = filtered_loans["home_ownership"].map({"MORTGAG":1,"RENT":2,"OTHER":3})
filtered_loans["emp_length"] = filtered_loans["emp_length"].replace({'years':'','other':'other'})
filtered_loans["emp_length"] = filtered_loans["emp_length"].apply(lambda x:int(x) if x!='other' else 0)
filtered_loans['verification_status'] = filtered_loans['verification_status'].map({"verified":1,"not_verified":2})
filtered_loans["term"] = filtered_loans["term"].map({" 36 months":0," 60 months":1})
print("Current shape of dataset :",filtered_loans.shape)
filtered_loans.head()
```

Current shape of dataset : (25744, 24)

Out[87]:

| | loan_amnt | term | installment | grade | emp_length | home_ownership | annual_inc | verification_status |
|---|-----------|------|-------------|-------|------------|----------------|------------|---------------------|
| 0 | 19800 | 0 | 666.00 | 5 | 7 | 3 | 78924.0 | |
| 1 | 35000 | 0 | 1177.27 | 5 | 10 | 2 | 95000.0 | |
| 2 | 20000 | 0 | 672.73 | 5 | 6 | 3 | 56000.0 | |
| 3 | 28000 | 1 | 635.37 | 5 | 5 | 3 | 96000.0 | |
| 4 | 23100 | 1 | 605.35 | 3 | 10 | 3 | 55000.0 | |

5 rows × 24 columns

Convert Nominal values to numeric values

```
In [88]: filtered_loans.to_csv("cleaned_loans_2015.csv",index=False)
```

```
In [89]: filtered_loans.shape
```

Out[89]: (25744, 24)

```
In [90]: X = filtered_loans.iloc[:,0:22]
y = filtered_loans.iloc[:,23]
```



```
In [91]: X.head()
```

```
Out[91]:
```

| | loan_amnt | term | installment | grade | emp_length | home_ownership | annual_inc | verification_statu |
|---|-----------|------|-------------|-------|------------|----------------|------------|--------------------|
| 0 | 19800 | 0 | 666.00 | 5 | 7 | 3 | 78924.0 | |
| 1 | 35000 | 0 | 1177.27 | 5 | 10 | 2 | 95000.0 | |
| 2 | 20000 | 0 | 672.73 | 5 | 6 | 3 | 56000.0 | |
| 3 | 28000 | 1 | 635.37 | 5 | 5 | 3 | 96000.0 | |
| 4 | 23100 | 1 | 605.35 | 3 | 10 | 3 | 55000.0 | |

5 rows × 22 columns

```
In [92]: y.head()
```

```
Out[92]: 0    1
         1    1
         2    1
         3    1
         4    1
         Name: loan_status, dtype: int64
```

```
In [93]: from sklearn import preprocessing,metrics
         from IPython.core.display import HTML

         filtered_loans.fillna(filtered_loans.mean(),inplace = True)
         HTML(filtered_loans.tail().to_html())
         print("Current shape of dataset :",filtered_loans.shape)
```

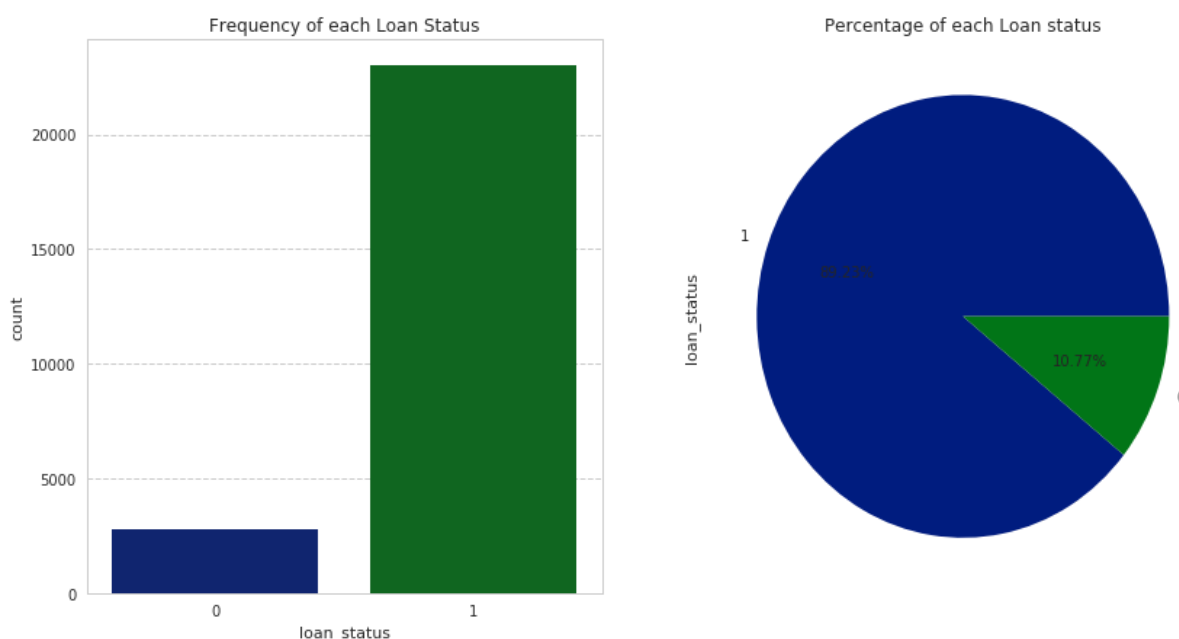
Current shape of dataset : (25744, 24)

```
In [166]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
plt.rcParams['figure.figsize'] = (12,8)

fig, axs = plt.subplots(1,2,figsize=(14,7))
sns.countplot(x='loan_status',data=loan_df,ax=axs[0])
axs[0].set_title("Frequency of each Loan Status")
loan_df.loan_status.value_counts().plot(x=None,y=None, kind='pie', ax=axs[1],auto
axs[1].set_title("Percentage of each Loan status")
plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu San
s

(prop.get_family(), self.defaultFamily[fonttext]))



```
In [94]: scl = preprocessing.StandardScaler() #instance of preprocessing
fields = filtered_loans.columns.values[:-1]
data_clean = pd.DataFrame(scl.fit_transform(filtered_loans[fields]), columns = fi
data_clean['loan_status'] = filtered_loans['loan_status']
data_clean['loan_status'].value_counts()
```

```
Out[94]: 1.0    22964
0.0     2767
Name: loan_status, dtype: int64
```

```
In [95]: from sklearn.utils import resample
loanstatus_0 = data_clean[data_clean["loan_status"]==0]
loanstatus_1 = data_clean[data_clean["loan_status"]==1]
loanstatus_0_upsampled = resample(loanstatus_0,
                                replace=True,      # sample with replacement
                                n_samples=20980,   # to match majority class
                                random_state=123) # reproducible results
data_clean = pd.concat([loanstatus_1, loanstatus_0_upsampled ])
data_clean = data_clean.sample(frac=1).reset_index(drop=True)
print("Current shape of dataset :",data_clean.shape)

loanstatus_0_upsampled.shape
```

Current shape of dataset : (43944, 24)

Out[95]: (20980, 24)

In [98]: `data_clean.corr()`

Out[98]:

| | loan_amnt | term | installment | grade | emp_length | home_owne |
|-----------------------------------|------------------|-------------|--------------------|--------------|-------------------|------------------|
| loan_amnt | 1.000000 | 0.455861 | 0.952071 | -0.225756 | 0.099306 | 0.12 |
| term | 0.455861 | 1.000000 | 0.217305 | -0.476183 | 0.064707 | 0.06 |
| installment | 0.952071 | 0.217305 | 1.000000 | -0.217058 | 0.084692 | 0.10 |
| grade | -0.225756 | -0.476183 | -0.217058 | 1.000000 | 0.019103 | 0.06 |
| emp_length | 0.099306 | 0.064707 | 0.084692 | 0.019103 | 1.000000 | 0.13 |
| home_ownership | 0.129773 | 0.068420 | 0.109254 | 0.063040 | 0.131661 | 1.00 |
| annual_inc | 0.317929 | 0.073557 | 0.306493 | 0.067256 | 0.087671 | 0.11 |
| verification_status | -0.216073 | -0.176614 | -0.196829 | 0.144564 | -0.015045 | 0.04 |
| dti | -0.017879 | 0.063224 | -0.015202 | -0.183153 | 0.013582 | -0.04 |
| delinq_2yrs | -0.011711 | -0.013231 | -0.005467 | -0.038706 | 0.017797 | 0.02 |
| inq_last_6mths | -0.020390 | -0.001253 | 0.011767 | -0.215075 | -0.001954 | -0.00 |
| open_acc | 0.175538 | 0.080293 | 0.163494 | -0.002158 | 0.035108 | 0.07 |
| pub_rec | -0.123417 | -0.071160 | -0.102815 | -0.047983 | -0.004923 | -0.00 |
| revol_bal | 0.316639 | 0.098604 | 0.304522 | 0.008651 | 0.072913 | 0.10 |
| revol_util | 0.117621 | 0.095156 | 0.120929 | -0.194006 | 0.052414 | 0.05 |
| total_acc | 0.184690 | 0.090346 | 0.163928 | 0.042440 | 0.110806 | 0.12 |
| initial_list_status | 0.082138 | 0.166243 | 0.011267 | 0.165295 | 0.035411 | 0.05 |
| last_pymnt_amnt | 0.725848 | 0.325575 | 0.679022 | -0.104151 | 0.101745 | 0.13 |
| collections_12_mths_ex_med | 0.000159 | 0.003060 | 0.003543 | -0.026857 | -0.003952 | -0.00 |
| acc_now_delinq | 0.024120 | 0.016880 | 0.023888 | -0.037193 | 0.005935 | -0.01 |
| tot_coll_amt | -0.026682 | -0.001487 | -0.025481 | -0.006535 | 0.018920 | 0.00 |
| tot_cur_bal | 0.332941 | 0.126096 | 0.305926 | 0.052522 | 0.108377 | 0.39 |
| total_rev_hi_lim | 0.352104 | 0.088110 | 0.327118 | 0.138139 | 0.070930 | 0.10 |
| loan_status | 0.010869 | -0.010601 | 0.012866 | 0.021460 | 0.006428 | -0.00 |

24 rows × 24 columns

```


In [99]: import numpy as np
from sklearn.svm import SVC
from sklearn.model_selection import learning_curve
from sklearn.model_selection import ShuffleSplit
from sklearn import linear_model, svm
from sklearn.metrics import average_precision_score
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(estimator, X, y, cv=cv,
                                                            n_jobs=n_jobs)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.grid()

    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1, color="g")
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")
    plt.legend(loc="best")
    return plt

X, y = data_clean.iloc[:, :-1].values, data_clean.iloc[:, -1].values
title = "Learning Curves (Logistic Regression)"
# Cross validation with 100 iterations to get smoother mean test and train
# score curves, each time with 20% data randomly selected as a validation set.
cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=0)
estimator = linear_model.LogisticRegression()
plot_learning_curve(estimator, title, X, y, ylim=(0.15, 0.90), cv=cv, n_jobs=4)
plt.show()

```



```
In [100]: import seaborn as sns
sns.set('talk', 'whitegrid', 'dark', font_scale=1, font='Ricty',rc={"lines.linewidth": 2})
def plotAUC(truth, pred, lab):
    fpr, tpr, _ = metrics.roc_curve(truth,pred)
    roc_auc = metrics.auc(fpr, tpr)
    lw = 2
    c = (np.random.rand(), np.random.rand(), np.random.rand())
    plt.plot(fpr, tpr, color= c,lw=lw, label= lab +' (AUC = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC curve') #Receiver Operating Characteristic
    plt.legend(loc="lower right")
```

```
In [101]: import itertools
from sklearn.metrics import confusion_matrix
def plot_confusion_matrix(model, normalize=False): # This function prints and plots the confusion matrix
    cm = confusion_matrix(y_test, model, labels=[0, 1])
    classes=["Will Pay", "Will Default"]
    cmap = plt.cm.Blues
    title = "Confusion Matrix"
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        cm = np.around(cm, decimals=3)
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

```
In [102]: X_train, X_test, y_train, y_test = train_test_split(data_clean.iloc[:, :-1], data_clean['y'],
bs_train, bs_test = train_test_split(data_clean, test_size = 0.2, random_state=42)
```

```
In [103]: #PCA (Principal Component Analysis)
from sklearn.decomposition import PCA
pca = PCA(n_components=10, whiten=True)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)
explained_variance = pca.explained_variance_ratio_
print('Expected Variance is ' + str(explained_variance))
```

```
Expected Variance is [ 0.17256429  0.08791379  0.07494317  0.06342185  0.057395
6  0.05214995
0.05023233  0.04913359  0.04449675  0.04376371]
```

```

In [104]: dataViz = data_clean
sns.set_context(context='notebook')
fig, ax = plt.subplots(figsize=(10,10))
corr = dataViz.corr()

# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.tril_indices_from(mask)] = True

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)


sns.heatmap(corr, cmap=cmap,linewidths=1, vmin=-1, vmax=1, square=True, cbar=True

```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu Sans

(prop.get_family(), self.defaultFamily[fonttext]))

Out[104]: <matplotlib.axes._subplots.AxesSubplot at 0x2a91df4c518>




```
In [105]: from sklearn.datasets import make_classification
from sklearn.ensemble import ExtraTreesClassifier

# Build a classification task using 3 informative features
X, y = make_classification(n_samples=6767,
                           n_features=19,
                           n_informative=10,
                           n_redundant=0,
                           n_repeated=0,
                           n_classes=2,
                           random_state=0,
                           shuffle=False)

# Build a forest and compute the feature importances
forest = ExtraTreesClassifier(n_estimators=250,
                              random_state=0)

forest.fit(X, y)
importances = forest.feature_importances_
std = np.std([tree.feature_importances_ for tree in forest.estimators_],
              axis=0)
indices = np.argsort(importances)[::-1]

# Print the feature ranking
print("Feature ranking:")

for f in range(X.shape[1]):
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))

# Plot the feature importances of the forest
plt.figure()
plt.title("Feature importances")
plt.bar(range(X.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")
plt.xticks(range(X.shape[1]), indices)
plt.xlim([-1, X.shape[1]])
plt.show()
```


Feature ranking:

1. feature 8 (0.134019)
2. feature 4 (0.099906)
3. feature 7 (0.098578)
4. feature 0 (0.095823)
5. feature 6 (0.077986)
6. feature 5 (0.066192)
7. feature 2 (0.065415)
8. feature 3 (0.063437)
9. feature 9 (0.059985)
10. feature 1 (0.044657)
11. feature 11 (0.022180)
12. feature 10 (0.021810)
13. feature 16 (0.021685)
14. feature 12 (0.021542)
15. feature 17 (0.021533)
16. feature 14 (0.021354)

- 17. feature 18 (0.021346)
- 18. feature 13 (0.021307)
- 19. feature 15 (0.021246)

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu San
s

(prop.get_family(), self.defaultFamily[fonttext]))

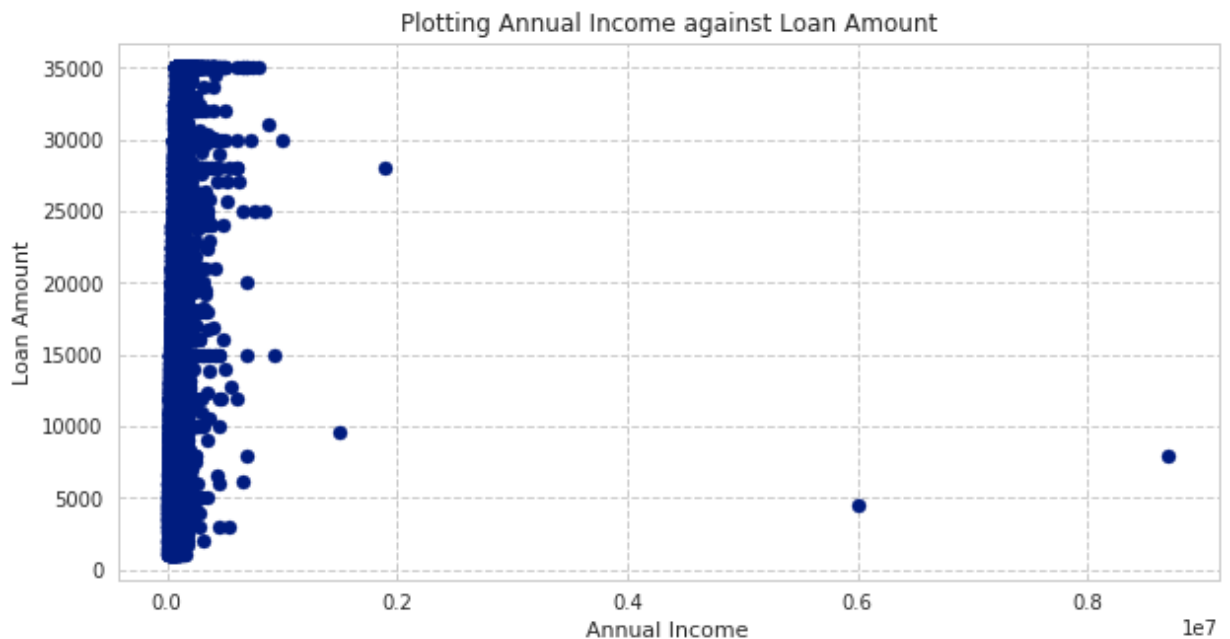


```
In [132]: plt.figure(figsize=(10,5))
plt.scatter(filtered_loans['annual_inc'], filtered_loans['loan_amnt'])
plt.title("Plotting Annual Income against Loan Amount")
plt.ylabel('Loan Amount')
plt.xlabel('Annual Income')
plt.show()

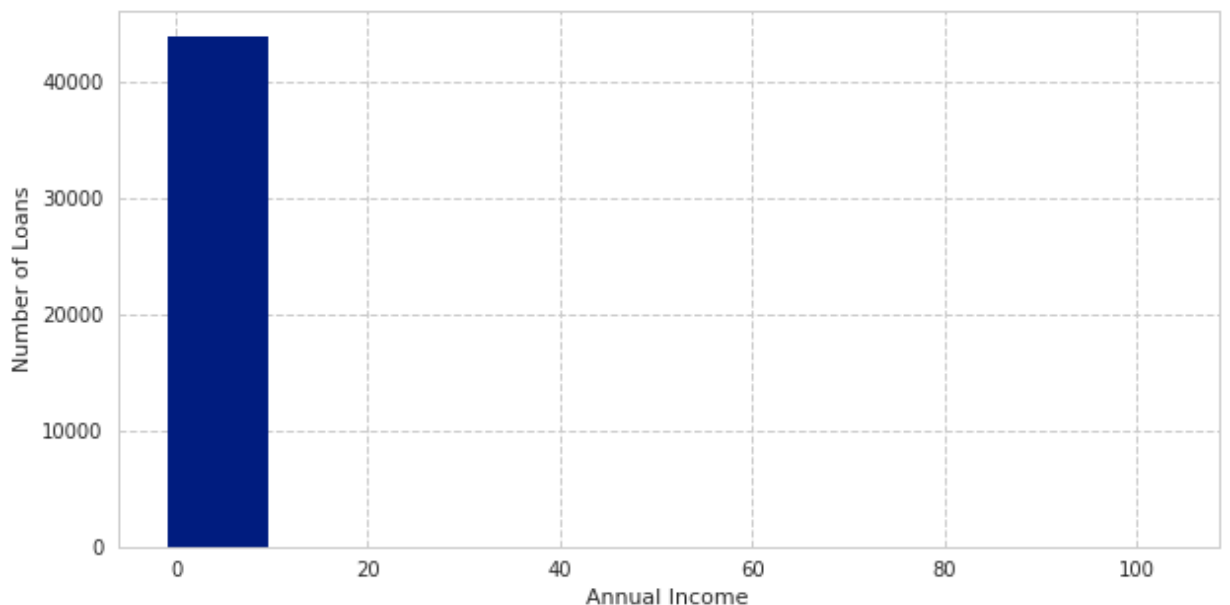
data_clean.annual_inc.hist(figsize=(10,5))
plt.ylabel('Number of Loans')
plt.xlabel('Annual Income')
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu Sans

(prop.get_family(), self.defaultFamily[fonttext]))



```
Out[132]: Text(0.5,0,'Annual Income')
```



```
In [133]: data_clean = data_clean[data_clean['annual_inc']<200000]
```

```
In [135]: X_Variables = ['emp_length', 'grade']
X = data_clean[X_Variables]
X = X.values
y = data_clean['loan_status'].values
clf = linear_model.LogisticRegression()
model = clf.fit(X,y)
model.score(X, y)
pd.DataFrame(list(zip(X_Variables,model.coef_.T)))
```

```
Out[135]:
```

| | 0 | 1 |
|--------------|-------------------|---|
| 0 emp_length | [0.0120505219261] | |
| 1 grade | [0.0426461226429] | |

Lets take a look at the co-efficient of grade a loan. For every additional increase in the grade "G" to "F" or in our case "1" to "2" the chance of the loan being paid off increases by 0.04 and it shows increase in emp_length by factor of 0.012

```
In [150]: from sklearn.model_selection import GridSearchCV
def cross_validation_best_parameters(model, param_grid):
    grid = GridSearchCV(model, param_grid,cv=10, scoring='accuracy')
    X=data_clean.iloc[:, :-1].values
    y=data_clean.iloc[:, -1].values
    grid.fit(X,y)
    mean_scores = [result.mean_validation_score for result in grid.grid_scores_]
    return mean_scores,grid.best_score_,grid.best_estimator_
logreg = linear_model.LogisticRegression(random_state=0)
c=[0.001, 0.01, 0.1, 1, 10, 100, 1000]
param_grid = dict(C=c)
mean_scores,Best_Accuracy, Best_classifier = cross_validation_best_parameters(log
print(Best_classifier)
```

```
LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=0, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)
```


```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:7
61: DeprecationWarning: The grid_scores_ attribute was deprecated in version 0.
18 in favor of the more elaborate cv_results_ attribute. The grid_scores_ attri
bute will not be available from 0.20
DeprecationWarning)
```

```
In [149]: clf_LR = linear_model.LogisticRegression(C=0.01)
clf_LR.fit(X_train,y_train)
LR_Predict = clf_LR.predict_proba(X_test)[:,-1]
LR_Predict_bin = clf_LR.predict(X_test)
LR_Accuracy = accuracy_score(y_test,LR_Predict.round())
print("Logistic regression accuracy is ",LR_Accuracy)
plt.figure(figsize=(6,6))
plot_confusion_matrix(LR_Predict_bin, normalize=True)
plt.show()
```

Logistic regression accuracy is 0.768094534712

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu Sans

(prop.get_family(), self.defaultFamily[fonttext]))



```

In [142]: from sklearn.neighbors import KNeighborsClassifier
          from sklearn.metrics import accuracy_score
          from sklearn.cross_validation import cross_val_score

myList = list(range(0,5))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train, y_train, cv=5, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()


```

The optimal number of neighbors is 3.

```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu San
s
(prop.get_family(), self.defaultFamily[fonttext]))

```

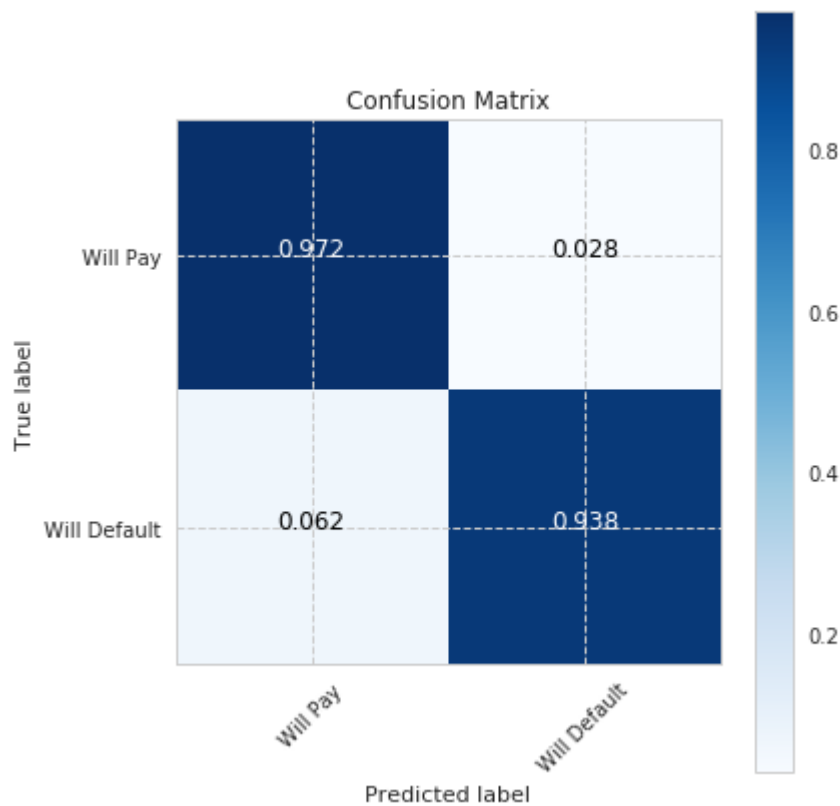


```
In [143]: from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=3)
neigh.fit(X_train,y_train)
LR_Predict_bin1 = neigh.predict(X_test)
LR_Accuracy = accuracy_score(y_test,LR_Predict_bin1.round())
print("KNN accuracy is ",LR_Accuracy)
plt.figure(figsize=(6,6))
plot_confusion_matrix(LR_Predict_bin1, normalize=True)
plt.show()
```

KNN accuracy is 0.954948301329

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu Sans

(prop.get_family(), self.defaultFamily[fonttext]))




```
In [147]: from sklearn import svm
clf_svm = svm.SVC()
clf_svm.fit(X_train,y_train)
#LR_Predict1 = clf_svm.predict_proba(X_test)[: ,1]
LR_Predict_bin2 = clf_svm.predict(X_test)
LR_Accuracy = accuracy_score(y_test,LR_Predict_bin2.round())
print("SVM accuracy is ",LR_Accuracy)


LR_Predict_train = clf_svm.predict(X_train)
LR_Accuracy1 = accuracy_score(y_train,LR_Predict_train.round())
print("SVM training accuracy is ",LR_Accuracy1)


plotAUC(y_test,LR_Predict_bin1,'K-NN')
plotAUC(y_test,LR_Predict_bin2,'SVM')
plt.show()
plt.figure(figsize=(6,6))
plot_confusion_matrix(LR_Predict_bin2, normalize=True)
plt.show()
```

SVM accuracy is 0.966765140325
 SVM training accuracy is 0.990578237576

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
 rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu San
 S

(prop.get_family(), self.defaultFamily[fonttext]))





```
In [151]: from sklearn.grid_search import RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
maxFeatures = range(1,data_clean.shape[1]-1)
param_dist = dict(max_features=maxFeatures)
rand = RandomizedSearchCV(rf, param_dist, cv=10, scoring='accuracy', n_iter=len(m
X=data_clean.iloc[:, :-1].values
y=data_clean.iloc[:, -1].values
rand.fit(X,y)
mean_scores = [result.mean_validation_score for result in rand.grid_scores_]
#print('Best Accuracy = '+str(rand.best_score_))
print(rand.best_estimator_)
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features=1, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
oob_score=False, random_state=None, verbose=0,
warm_start=False)
```

```
In [165]: randomForest = RandomForestClassifier(bootstrap=True,criterion = "gini",max_featu


randomForest.fit(X_train,y_train)
rfPredict = randomForest.predict(X_test)
rfPredicttrain = randomForest.predict(X_train)
rfAccuracytrain = accuracy_score(y_train,rfPredicttrain)
rfPredictproba = randomForest.predict_proba(X_test)[:,-1] #for ROC curve
rfAccuracy = accuracy_score(y_test,rfPredict)
roc_score = metrics.roc_auc_score(y_test,rfPredict)
print(rfAccuracy)
print(rfAccuracytrain)
plt.figure(figsize=(6,6))
plot_confusion_matrix(rfPredict, normalize=True)
plt.show()
```

0.84047267356

0.999630519121

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\font_manager.py:1316: Use
rWarning: findfont: Font family ['Ricty'] not found. Falling back to DejaVu Sans

(prop.get_family(), self.defaultFamily[fonttext]))



```

In [140]: from sklearn.model_selection import cross_val_score
import numpy as np
X=data_clean.iloc[:, :-1].values
y=data_clean.iloc[:, -1].values
clf = RandomForestClassifier()
print(np.mean(cross_val_score(clf, X_train, y_train, cv=15)))

param_grid = {
    'n_estimators': [5, 10, 15, 20],
    'max_depth': [2, 5, 6, 7, 8, 9]
}
from sklearn.grid_search import GridSearchCV

grid_clf = GridSearchCV(clf, param_grid, cv=15)
grid_clf.fit(X_train, y_train)
grid_clf.best_estimator_
grid_clf.best_params_
grid_clf.grid_scores_

```

0.906330308745

```

Out[140]: [mean: 0.72178, std: 0.04284, params: {'max_depth': 2, 'n_estimators': 5},
mean: 0.75559, std: 0.02661, params: {'max_depth': 2, 'n_estimators': 10},
mean: 0.76575, std: 0.02620, params: {'max_depth': 2, 'n_estimators': 15},
mean: 0.77813, std: 0.02065, params: {'max_depth': 2, 'n_estimators': 20},
mean: 0.81877, std: 0.02123, params: {'max_depth': 5, 'n_estimators': 5},
mean: 0.83244, std: 0.01893, params: {'max_depth': 5, 'n_estimators': 10},
mean: 0.84075, std: 0.01509, params: {'max_depth': 5, 'n_estimators': 15},
mean: 0.84463, std: 0.01556, params: {'max_depth': 5, 'n_estimators': 20},
mean: 0.84075, std: 0.01172, params: {'max_depth': 6, 'n_estimators': 5},
mean: 0.85849, std: 0.01110, params: {'max_depth': 6, 'n_estimators': 10},
mean: 0.85960, std: 0.01494, params: {'max_depth': 6, 'n_estimators': 15},
mean: 0.86717, std: 0.01290, params: {'max_depth': 6, 'n_estimators': 20},
mean: 0.86255, std: 0.01231, params: {'max_depth': 7, 'n_estimators': 5},
mean: 0.87087, std: 0.01620, params: {'max_depth': 7, 'n_estimators': 10},
mean: 0.87604, std: 0.01555, params: {'max_depth': 7, 'n_estimators': 15},
mean: 0.88417, std: 0.01753, params: {'max_depth': 7, 'n_estimators': 20},
mean: 0.87068, std: 0.01499, params: {'max_depth': 8, 'n_estimators': 5},
mean: 0.88177, std: 0.01700, params: {'max_depth': 8, 'n_estimators': 10},
mean: 0.89174, std: 0.01349, params: {'max_depth': 8, 'n_estimators': 15},
mean: 0.89248, std: 0.01620, params: {'max_depth': 8, 'n_estimators': 20},
mean: 0.87308, std: 0.01432, params: {'max_depth': 9, 'n_estimators': 5},
mean: 0.89710, std: 0.01639, params: {'max_depth': 9, 'n_estimators': 10},
mean: 0.89987, std: 0.01253, params: {'max_depth': 9, 'n_estimators': 15},
mean: 0.90227, std: 0.01516, params: {'max_depth': 9, 'n_estimators': 20}]

```